

# MitoChime: A Machine Learning Pipeline for Detecting PCR-Induced Chimeras in Mitochondrial Illumina Reads

Duran, Lin, Pailden

University of the Philippines Visayas

December 9, 2025

# Outline

- 1 Introduction
- 2 Problem Statement & Proposed Solution
- 3 Objectives
- 4 Scope and Limitations
- 5 Methodology

## Next Generation Sequencing (NGS)

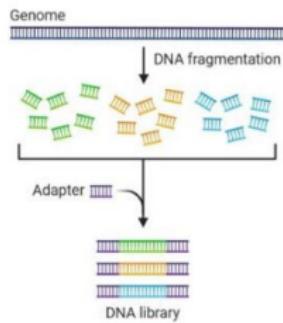


*Source: University of the Philippines  
Visayas, 2022*

## Illumina Seq Workflow

### Step 1. Library Preparation

#### ① Library preparation



Source: Microbe Notes, 2024

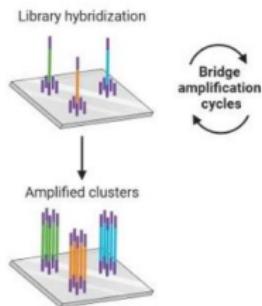


Source: Philippine Genome Center Visayas, 2025

## Illumina Seq Workflow

### Step 2. Library Bridge Amplification (PCR)

#### ② DNA library bridge amplification



Source: Microbe Notes, 2024

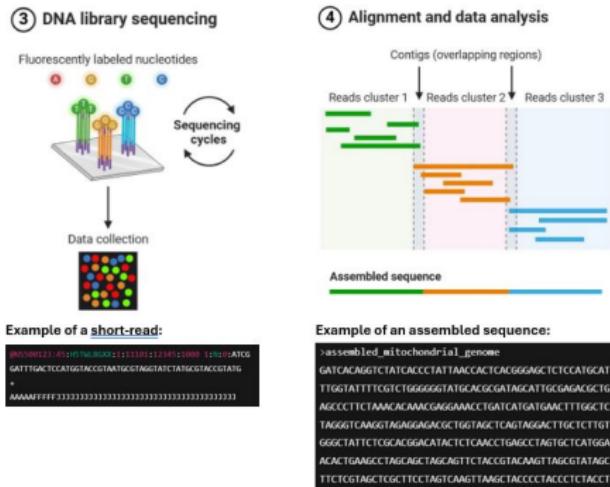


Source: Philippine Genome Center Visayas, 2025

## Real-World Problem

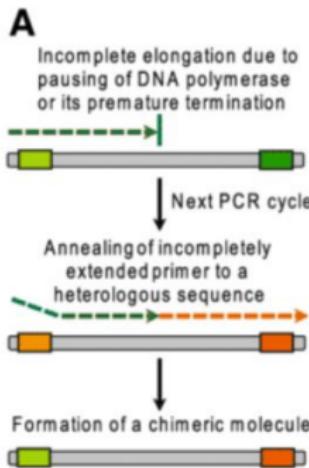
# Illumina Seq Workflow

### Step 3. Sequencing and Alignment



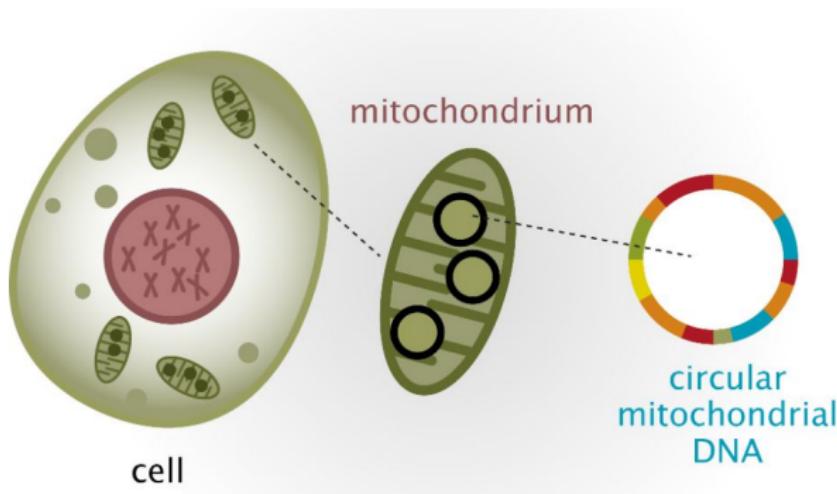
# Real-World Problem

## PCR-Chimera Formation



Source: Omelina et al., 2024

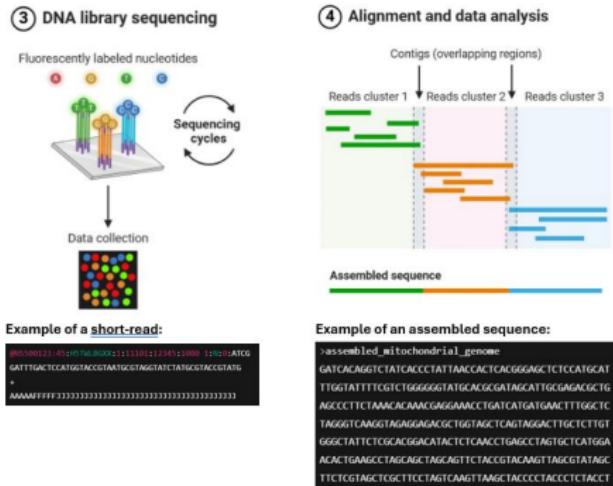
## The Mitochondrial Genome



Source: UZ Bruzel., 2020

## Real-World Problem

## Disrupts Genome Assembly



# Existing Approaches

Table 2.1: Comparison of Chimera Detection Approaches and Tools

Method / Tool	Core Approach	Key Limitations
Reference-based Detection	Compares each query sequence against curated databases of verified, non-chimeric sequences; evaluates segment similarity to identify mosaic patterns.	Accuracy depends on database completeness; performs poorly for novel taxa or missing parents; limited sensitivity for low-divergence chimeras.
De novo Detection	Identifies chimeras using only internal dataset structure; leverages abundance hierarchy and compositional similarity to infer whether low-abundance sequences can be reconstructed from abundant parents.	Assumes true sequences are more abundant; fails when amplification bias distorts abundances; struggles when parental sequences are similarly abundant or highly similar.
UCHIME	Alignment-based model that partitions the query into segments, identifies parent candidates, and computes a chimera score via a three-way alignment; supports reference and de novo modes.	Reduced accuracy for very closely related parents (<0.8% divergence); sensitive to incomplete databases; de novo mode fails if parents are absent or not sufficiently more abundant.
UCHIME2	Updated UCHIME with improved benchmarking (CHSIMA) and multiple sensitivity/specificity presets; better handles incomplete references and dataset variability.	"Fake models" limit theoretical accuracy; genuine variants may mimic chimeras; not recommended as a standalone step in OTU or denoising pipelines due to increased false positives/negatives.
CATCh	First ensemble ML model for 16S chimera detection; integrates outputs of UCHIME, ChimeraSlayer, DECIPHER, Pintail, and Perseus using an SVM to boost overall prediction accuracy.	Performance constrained by underlying tools; ML model cannot capture features not present in component algorithms; may misclassify in highly novel or low-coverage datasets.
ChimPipe	Pipeline for detecting biological chimeras in RNA-seq using discordant paired-end reads and split-read alignments; identifies isoforms and breakpoint coordinates.	Requires high-quality genome and annotation; tailored to RNA-seq rather than amplicons; computationally intensive; limited to organisms with available reference genomes.

# Problem Statement & Proposed Solution

- **Problem Statement:** Chimeric sequencing reads can disrupt mitochondrial genome assembly, but current assembly pipelines assume artifact-free input and existing chimera detection tools are not designed specifically for organellar, particularly mitochondrial datasets, leaving assemblies vulnerable to undetected artifacts.
- **Proposed Solution:** A machine-learning pipeline designed to detect PCR-induced chimeric reads using both alignment-based and sequence-derived features to improve the quality and reliability of downstream mitochondrial genome assemblies.

# General Objective

- Develop and evaluate a machine-learning pipeline (MitoChime) to detect PCR-induced chimeric reads in *S. lemuru* mitochondrial sequencing data to improve downstream assembly quality.

# Specific Objectives

- ① Construct simulated *Sardinella lemuru* Illumina paired-end datasets containing both clean and PCR-induced chimeric reads.
- ② Extract alignment-based and sequence-based features such as k-mer composition, microhomology, and split-alignment counts from both clean and chimeric reads
- ③ Train, validate, and compare supervised machine learning models for classifying reads as clean or chimeric.
- ④ Determine feature importance and identify indicators of PCR-induced chimerism.
- ⑤ Integrate the optimized classifier into a modular and interpretable pipeline deployable on standard computing environments at PGC Visayas.

# Scope of the Study

- Focuses on PCR-induced chimeric reads in *Sardinella lemuru* mitochondrial sequencing data to:
  - to limit interspecific variation in mitochondrial genome size, GC content, and repetitive regions so that differences in read patterns can be attributed more directly to PCR-induced chimerism
  - to align the analysis with relevant *S. lemuru* sequencing projects at PGC Visayas
  - to take advantage of the availability of *S. lemuru* mitochondrial assemblies and raw datasets in public repositories such as the National Center for Biotechnology Information (NCBI), which facilitates reference selection and benchmarking
  - to develop a tool that directly supports local studies on *S. lemuru* population structure and fisheries management produce tools applicable to local population and fisheries studies

# Scope of the Study

- Uses wgsim-based simulations and selected empirical mitochondrial datasets
- Analysis targets low-dimensional alignment and sequence features (k-mers, GC content, clipping, split alignments) to maintain interpretability and computational accessibility
- Long-read platforms and other taxa are not included

# Key Exclusions

- Naturally occurring chimeras
- NUMTs
- Large-scale nuclear genome rearrangements
- High-dimensional deep learning embeddings

## Other Limitations

- No simulations with variable sequencing error rates
- No testing of alternative parameter settings (k-mer length, microhomology windows)
- Reliance on supervised machine learning may limit detection of novel/unknown chimeric patterns

# Methodology

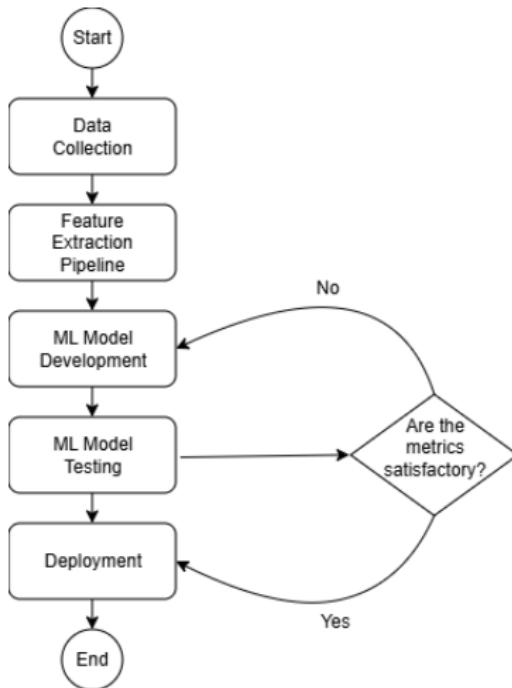


Figure: Process Diagram of the Special Project

# Data Collection

The *S. lemuru* mitochondrial reference genome (NCBI: NC\_039553.1) was downloaded in FASTA format and used as the basis for generating simulated reads.

# Data Preprocessing

- A Python script was used to generate the reads.
- Clean reads were produced with wgsim from the reference genome.
- A chimeric reference was created by creating a custom script to combine non-adjacent segments with microhomology
- Chimeric reads were simulated with wgsim.
- All reads were mapped with minimap2 to extract alignment information.
- SAM/BAM files were converted, sorted, and indexed with samtools.

# Data Preprocessing

- Final dataset: 40k reads, roughly balanced between clean and chimeric (19,984 clean reads and 20,000 chimeric).
- Some of the clean reads failed to align due to the set error rate.

# Data Preprocessing

```
NC_039553_1_3_540_8:0:0 6:0:0 ef2      163    NC_039553.1   3       60      150M =      391    538
    TGGTAGCTTAAACAAAGCATAACACTGAAAGATGGCTCGTATAAGCCCCAACGCACGTGAAAGTTTGCTGCTGCTTATTATCAGCTTACCGGAATTACACCAGCGAGCCCTCCGCCGCGCCGTGAGGATGCCCTCA
    .....
AS:i:220    nn:i:0  tp:A:P  cm:i:8  si:i:164    s2:i:0  def:f:0.0533  r1:i:0
NC_039553_1_4_430_13:0:0 11:0:0_243d  163    NC_039553.1   4       60      150M =      281    427
    GGTTAGCTTAAACAAAGCATAACACTGCAAGATGATCCGCTGGCCGTGATAAGCCGAGCAGGAGTGAAGTTTGCTGAGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCCGCCGCCCCGTGAGGATGCCCTCAG
    .....
AS:i:170    nn:i:0  tp:A:P  cm:i:9  si:i:135    s2:i:0  def:f:0.0867  r1:i:0
NC_039553_1_5_495_6:0:0 11:0:0_1d49  163    NC_039553.1   5       60      150M =      346    491
    GTGTAGCTTACACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTGATCAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGTTTCCCCAATTACACATGCGAGCCTCCGCCGCCCCGTGAGGATGCCCTCAGC
    .....
AS:i:240    nn:i:0  tp:A:P  cm:i:12 si:i:148    s2:i:0  def:f:0.04   r1:i:0
NC_039553_1_6_523_6:0:0 9:0:0_82c   163    NC_039553.1   6       60      150M =      374    518
    TGTAGCTTAAACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTGATAAGCCCCAACAGCTGAAAGTTAGGTCTGGCTATTACAGCTTACCCAAATTACACATGCGCCCTCCGCCGCCCCGTGAGGATGCCCTCAGCC
    .....
AS:i:240    nn:i:0  tp:A:P  cm:i:10 si:i:157    s2:i:0  def:f:0.04   r1:i:0
NC_039553_1_9_574_7:0:0 7:0:0_181b  163    NC_039553.1   9       60      150M =      425    566
    AGCTTAAACAAAGCATAACACTGCAAGATGTTAAGCTGGCCGTGATAAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCCGCCGCCCCGTGAGGCTGCCCTCCGCCCTCC
    .....
AS:i:230    nn:i:0  tp:A:P  cm:i:12 si:i:176    s2:i:0  def:f:0.0467  r1:i:0
NC_039553_1_10_391_9:0:0 8:0:0_256b  99     NC_039553.1   10      60      150M =      242    382
    GCTTAAACAAAGCTAACACTGAAAGATGTTAAGATGGCCGTGATAAGCCCCAACAGCACAGAAAGTTAGGTCTGGCTTATTACAGCTTACCCAAATTAGACATGCGAGCCTCCGCCGCCCCGTGAGTGTGCCCTCAGCTCC
    .....
AS:i:210    nn:i:0  tp:A:P  cm:i:5  si:i:156    s2:i:0  def:f:0.06   r1:i:0
NC_039553_1_11_509_6:0:0 11:0:0_a19  99     NC_039553.1   11      60      150M =      360    499
    CCTCACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTATAAGCCCACAGCACTGAAAGTTAGGTCTGGCTTATTATGAGCTTACCCAAATTACACATGCGATCCGCCGCCCCGTGAGGATGCCCTCAGCTCCCG
    .....
AS:i:242    nn:i:0  tp:A:P  cm:i:10 si:i:150    s2:i:0  def:f:0.04   r1:i:0
NC_039553_1_12_427_9:0:0 9:0:0_157   163    NC_039553.1   12      60      150M =      278    416
    TTAAACAAGCATAACACTGAAAGATTCAAGATGGCCGTATAAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCGGGGGCCCCGTGAGGATGCCCTCCGCCCTCCGT
    .....
AS:i:210    nn:i:0  tp:A:P  cm:i:8  si:i:158    s2:i:0  def:f:0.06   r1:i:0
```

Figure: SAM File of Clean Reads

## Data Preprocessing

## Figure: SAM File of Chimeric Reads

# Feature Extraction Pipeline

- BAM files were processed with a Python script (`extract_features.py`) to build a TSV feature matrix.
- Used Pysam for parsing alignments and NumPy for computation.

# Feature Extraction Pipeline

- Focused on three features linked to PCR-induced chimeras:
  - ① **Supplementary Alignment (SA)**: Detects split alignments; counts and metrics extracted from SA tags
  - ② **K-mer Composition Difference**: Breakpoints inferred; left/right segments compared using cosine and JS metrics.
  - ③ **Microhomology**: Overlap at junction quantified (length + GC content) within a defined window.
- Pipeline design and outputs to be validated by experts.

# Feature Extraction Pipeline

read_id	label	read_length	mean_basenr	ref_start	3'strand	mapp	cigar	has_ss	ss_sa_count	num_segs_in_ss	sa_diff	co_sa_min	dk_sa_max	d_sa_mean	sa_same	sa_opp	st_sa_max	nr_sa_mean	sa_min	nr_sa_mean	softclip_le	softclip_le	n_total	clipp	breakpoint	kmer	cosi	kmer_js_d	microhom	microhom
NC_09950	0	150	13 NC_09950	3	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.9720	0.97143	1	0		
NC_09950	0	150	13 NC_09950	4	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98591	0.98571	1	0		
NC_09950	0	150	13 NC_09950	5	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.95887	0.95714	0	0		
NC_09950	0	150	13 NC_09950	6	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97183	0.97143	1	1		
NC_09950	0	150	13 NC_09950	7	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98644	0.98571	1	0		
NC_09950	0	150	13 NC_09950	10	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97296	0.97143	0	0		
NC_09950	0	150	13 NC_09950	11	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97296	0.97143	1	0		
NC_09950	0	150	13 NC_09950	12	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	1	1		
NC_09950	0	150	13 NC_09950	12	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98648	0.98571	1	1		
NC_09950	0	150	13 NC_09950	12	0	24	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.95680	0.95714	1	1		
NC_09950	0	150	13 NC_09950	14	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	0	0		
NC_09950	0	150	13 NC_09950	15	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98649	0.98571	1	0		
NC_09950	0	150	13 NC_09950	17	0	60	146M4S	0	0	1	0	0	0	0	0	0	0	0	0	0	4	4	146	0	0.5	0	0			
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98649	0.98571	3	0		
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	3	0		
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98629	0.98571	3	0		
NC_09950	0	150	13 NC_09950	19	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	3	0		
NC_09950	0	150	13 NC_09950	20	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	0	0		
NC_09950	0	150	13 NC_09950	21	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0		
NC_09950	0	150	13 NC_09950	23	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98667	0.98571	0	0		
NC_09950	0	150	13 NC_09950	25	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98629	0.98571	0	0		
NC_09950	0	150	13 NC_09950	26	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0		
NC_09950	0	150	13 NC_09950	32	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97238	0.97143	2	1		
NC_09950	0	150	13 NC_09950	34	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	0	0		
NC_09950	0	150	13 NC_09950	34	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	0	0		
NC_09950	0	150	13 NC_09950	35	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0		
NC_09950	0	150	13 NC_09950	36	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98648	0.98571	0	0		
NC_09950	0	150	13 NC_09950	38	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0		
NC_09950	0	150	13 NC_09950	39	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98684	0.98571	0	0		
NC_09950	0	150	13 NC_09950	41	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97296	0.97143	2	0.5		
NC_09950	0	150	13 NC_09950	43	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	0	0		

Figure: TSV Dataset showing Clean Reads

# Feature Extraction Pipeline

1	read_id	label	-3	read_id	mean	ref_nuc	ref_stal	strand	w	mpg	cigar	has_sai	sai_csq	num_sai	sai_dif	w	sai_min	w	sai_max	w	sai_pct	w	sai_max	w	sai_min	w	sai_max	w	sai_pct	w	sai_pct	w	total_c	w	breakup	w	kmer_c	w	kmer_j	w	microf	w	microf	w	biology_gc
19985	chimeric_1	1	150	40 NC_03956	40	1	60 150M	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0					
19986	chimeric_1	1	150	40 NC_03956	53	0	0 000000	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0							
19987	chimeric_1	1	150	40 NC_03956	65	0	0 000000	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
19988	chimeric_1	1	150	40 NC_03956	66	0	0 000000	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
19989	chimeric_1	1	150	40 NC_03956	67	0	0 000000	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
19990	chimeric_1	1	150	40 NC_03956	67	1	60 110M32S	1	1	2	0	4246	4246	4246	0	1	10	10	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0							
19991	chimeric_1	1	150	40 NC_03956	69	1	60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
19992	chimeric_1	1	150	40 NC_03956	76	0	0 60 100M41S	1	1	2	0	4237	4237	4237	0	1	16	16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0							
19993	chimeric_1	1	150	40 NC_03956	77	1	60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
19994	chimeric_1	1	150	40 NC_03956	79	0	0 60 100M44S	1	1	2	0	4234	4234	4234	0	1	17	17	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0							
19995	chimeric_1	1	150	40 NC_03956	84	0	0 60 120M38S	1	1	2	0	5197	5197	5197	0	1	10	10	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0							
19996	chimeric_1	1	150	40 NC_03956	85	0	0 60 110M39S	1	1	2	0	5180	5180	5180	0	1	20	20	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
19997	chimeric_1	1	150	40 NC_03956	88	0	0 60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
19998	chimeric_1	1	150	40 NC_03956	89	0	0 60 150M39M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
19999	chimeric_1	1	150	40 NC_03956	89	0	0 60 305120M	1	1	2	0	1973	1973	1973	0	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20000	chimeric_1	1	150	40 NC_03956	89	0	0 60 415109M	1	1	2	0	1962	1962	1962	0	1	15	15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20001	chimeric_1	1	150	40 NC_03956	89	1	60 96M54S	1	1	2	0	4224	4224	4224	0	1	48	48	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20002	chimeric_1	1	150	40 NC_03956	89	1	60 355117M	1	1	2	0	1970	1970	1970	0	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20003	chimeric_1	1	150	40 NC_03956	90	0	0 60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20004	chimeric_1	1	150	40 NC_03956	90	1	60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20005	chimeric_1	1	150	40 NC_03956	90	1	60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20006	chimeric_1	1	150	40 NC_03956	91	0	0 60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0										
20007	chimeric_1	1	150	40 NC_03956	91	0	0 60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0										
20008	chimeric_1	1	150	40 NC_03956	91	1	60 150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20009	chimeric_1	1	150	40 NC_03956	91	1	60 94M65S	1	1	2	0	4222	4222	4222	0	1	52	52	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20010	chimeric_1	1	150	40 NC_03956	92	0	0 265121M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20011	chimeric_1	1	150	40 NC_03956	92	0	0 275278M	1	1	2	0	3064	3064	3064	0	1	60	60	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0								
20012	chimeric_1	1	150	40 NC_03956	92	0	0 665848M	1	1	2	0	3070	3070	3070	0	1	59	59	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									
20013	chimeric_1	1	150	40 NC_03956	92	0	0 665849M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0										
20014	chimeric_1	1	150	40 NC_03956	92	0	0 665124M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0										
20015	chimeric_1	1	150	40 NC_03956	92	0	0 653147M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0										
20016	chimeric_1	1	150	40 NC_03956	92	0	0 545994M	1	1	2	0	3082	3082	3082	0	1	30	30	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0									

Figure: TSV Dataset showing Chimeric Reads

# Dataset construction and split

- Simulated feature tables:
  - Clean reads (label 0)
  - PCR-induced chimeras (label 1)
- `build_datasets.py`:
  - Concatenate tables
  - Shuffle rows (avoid file-order artefacts)
- 80/20 **stratified** train–test split
- Test set held out and used **only once** at the end

# Validation strategy

- Layer 1: 80/20 stratified train–test split
- Layer 2: 5-fold stratified cross-validation on training set
  - Train on 4 folds, validate on 1
  - Rotate so each fold is validation once
- Layer 3: Final evaluation on held-out test set
- Hyperparameter tuning:
  - RandomizedSearchCV inside CV for top models
- Goal: stable estimates and **unbiased** test performance

# Model zoo and preprocessing pipeline

- **Baseline:** dummy majority-class classifier
- **Linear models:** logistic regression, calibrated linear SVM
- **Tree ensembles:**
  - Random Forest, Extra Trees
  - Gradient Boosting, XGBoost, LightGBM, CatBoost
- **Others:** bagging trees, k-NN, Gaussian NB, shallow MLP
- Common scikit-learn pipeline:
  - Median imputation (numeric missing values)
  - Standardisation (zero mean, unit variance)
- Ensures a **fair comparison** across models

# Effect of hyperparameter tuning

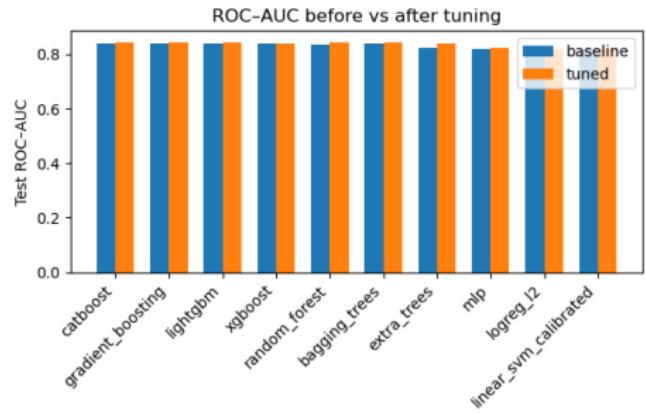
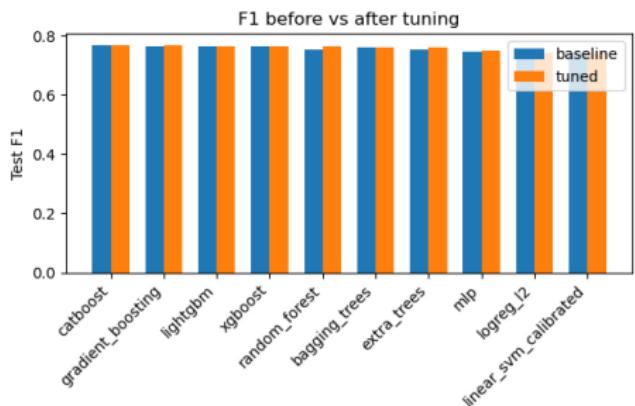


Figure: Test F1: baseline vs tuned.

Figure: Test ROC-AUC: baseline vs tuned.

- Tuning done with `RandomizedSearchCV` on training set
- Small but consistent gains ( $\Delta F1, \Delta AUC \approx 0.001\text{--}0.01$ )
- Top-ranked models remain the same (CatBoost, Gradient Boosting, LightGBM)

# ROC and precision–recall curves

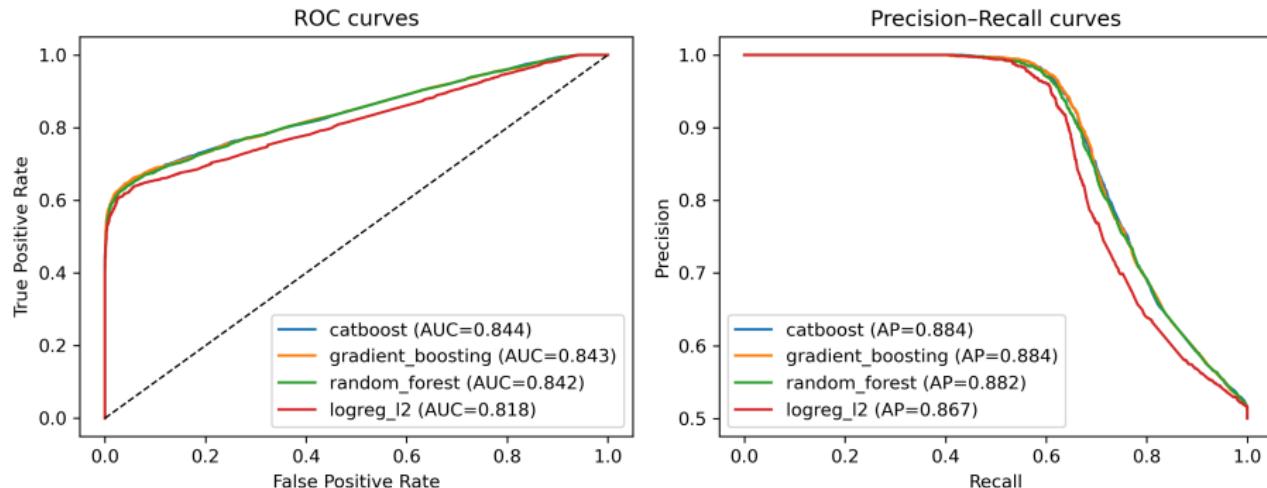
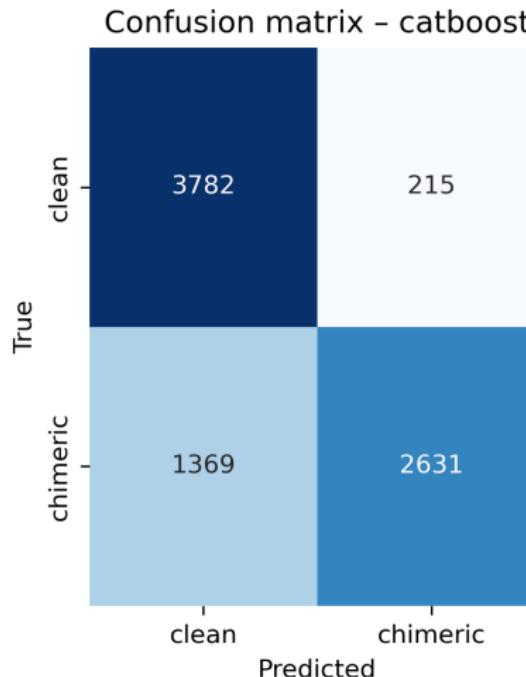


Figure: ROC (left) and PR (right) curves for CatBoost, Gradient Boosting, Random Forest, and logistic regression.

- Ensembles: ROC–AUC  $\approx 0.84$ ; logreg:  $\approx 0.82$
- Average precision  $\approx 0.88$  for ensembles
- Precision  $> 0.9$  up to recall  $\approx 0.5\text{--}0.6$

# Confusion matrix: CatBoost (test set)



- Clean reads:
  - Recall  $\approx 0.95$  ( $3782 / 3997$ )
- Chimeric reads:
  - Precision  $\approx 0.92$
  - Recall  $\approx 0.66$  ( $2631 / 4000$ )
- Behaviour at default threshold:
  - **Conservative chimera filter**
  - Protects clean reads, misses some subtle chimeras

Figure: Confusion matrix heatmap for CatBoost.

# Top features for CatBoost

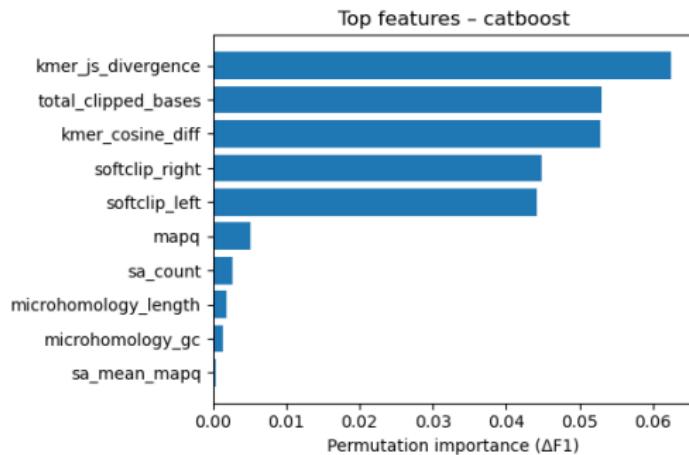


Figure: Permutation importance ( $\Delta F1$ ) for CatBoost.

- Strongest signals:
  - kmer\_js\_divergence
  - total\_clipped\_bases
  - kmer\_cosine\_diff
- Also important:
  - Left/right soft-clipping
  - Mapping quality (MAPQ)
  - SA count (supplementary alignments)
- Consistent with PCR chimera junctions

# Feature family importance

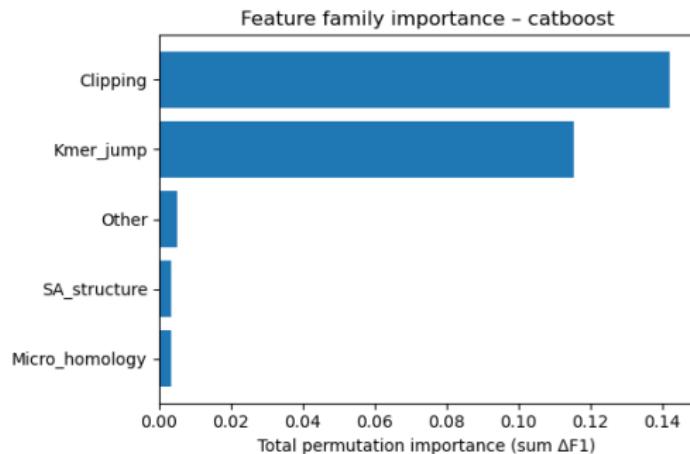


Figure: Aggregated feature families for CatBoost.

- Aggregated permutation importance:
  - **Clipping** features dominate
  - **K-mer jump** features also strong
- Smaller contributions:
  - SA structure
  - Micro-homology
  - Other alignment context
- Same pattern for Gradient Boosting and Random Forest

# Summary

- Tree-based ensembles (CatBoost, Gradient Boosting, LightGBM) consistently outperform linear baselines.
- Best models reach  $F1 \approx 0.77$  and ROC–AUC  $\approx 0.84$  on held-out reads.
- Key predictive signals match chimera biology:
  - k-mer composition jumps along the read
  - extensive soft-clipping and total clipped bases
- At the default threshold, the filter is conservative:
  - preserves most clean reads
  - removes a substantial fraction of chimeras
- Overall: our feature set and model ensemble provide a practical pre-filter before mitochondrial assembly.