

Machine Learning Pipeline for Detecting PCR-Induced Chimeric Reads

MitoChime: Organellar Chimera Detection from Per-Read Features

Duran, Lin, Pailden

University of the Philippines Visayas
Philippine Genome Center Visayas

December 8, 2025

Outline

- 1 Objectives
- 2 Scope and Limitations
- 3 Methodology
- 4 Train–Test Split and Validation
- 5 Model Zoo and Training
- 6 Metrics and Interpretation
- 7 Results
- 8 Discussion and Conclusion

General Objective

- Develop and evaluate a machine-learning pipeline (MitoChime) to detect PCR-induced chimeric reads in *S. lemuru* mitochondrial sequencing data to improve downstream assembly quality.

Specific Objectives

- ① Construct simulated *Sardinella lemuru* Illumina paired-end datasets containing both clean and PCR-induced chimeric reads.
- ② Extract alignment-based and sequence-based features such as k-mer composition, junction complexity, and split-alignment counts from both clean and chimeric reads
- ③ Train, validate, and compare supervised machine learning models for classifying reads as clean or chimeric.
- ④ Determine feature importance and identify indicators of PCR-induced chimerism.
- ⑤ Integrate the optimized classifier into a modular and interpretable pipeline deployable on standard computing environments at PGC Visayas.

Scope of the Study

- Focuses on PCR-induced chimeric reads in *Sardinella lemuru* mitochondrial sequencing data to:
 - to limit interspecific variation in mitochondrial genome size, GC content, and repetitive regions so that differences in read patterns can be attributed more directly to PCR-induced chimerism
 - to align the analysis with relevant *S. lemuru* sequencing projects at PGC Visayas
 - to take advantage of the availability of *S. lemuru* mitochondrial assemblies and raw datasets in public repositories such as the National Center for Biotechnology Information (NCBI), which facilitates reference selection and benchmarking
 - to develop a tool that directly supports local studies on *S. lemuru* population structure and fisheries management produce tools applicable to local population and fisheries studies

Scope of the Study

- Uses wgsim-based simulations and selected empirical mitochondrial datasets
- Analysis targets low-dimensional alignment and sequence features (k-mers, GC content, clipping, split alignments) to maintain interpretability and computational accessibility
- Long-read platforms (Nanopore, PacBio) and other taxa are not included

Key Exclusions

- Naturally occurring chimeras
- NUMTs
- Large-scale nuclear genome rearrangements
- High-dimensional deep learning embeddings

Other Limitations

- No simulations with variable sequencing error rates
- No testing of alternative parameter settings (k-mer length, microhomology windows)
- Reliance on supervised machine learning may limit detection of novel/unknown chimeric patterns

Methodology

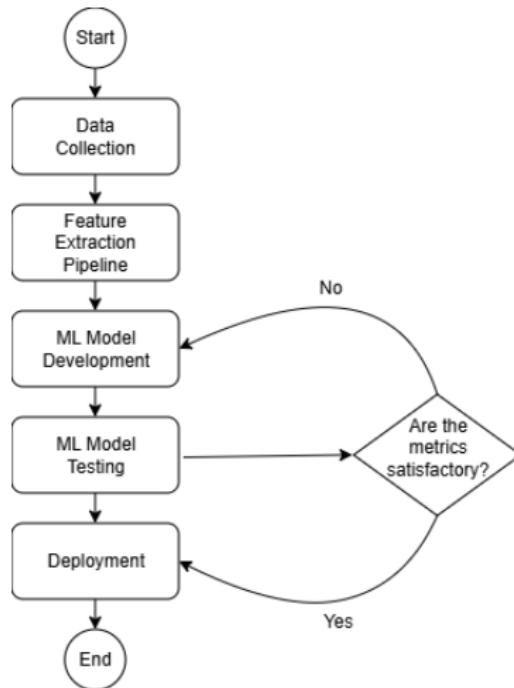


Figure: Process Diagram of the Special Project

Data Collection

The *S. lemuru* mitochondrial reference genome (NCBI: NC_039553.1) was downloaded in FASTA format and used as the basis for generating simulated reads.

Data Preprocessing

- A Python script was used to generate the reads.
- Clean reads were produced with wgsim from the reference genome.
- A chimeric reference was created by creating a custom script to combine non-adjacent segments with microhomology
- Chimeric reads were simulated with wgsim.
- All reads were mapped with minimap2 to extract alignment information.
- SAM/BAM files were converted, sorted, and indexed with samtools.

Data Preprocessing

- Final dataset: 40k reads, roughly balanced between clean and chimeric (19,984 clean reads and 20,000 chimeric).
- Some of the clean reads failed to align due to the set error rate.

Data Preprocessing

```
NC_039553_1_3_540_8:0:0 6:0:0 ef2      163    NC_039553.1   3       60     150M =      391    538
    TGGTAGCTTAAACAAAGCATAACACTGAAAGATGGCTCGTATAAGCCCCAACGCACGTGAAAGTTTGCTGCTGCTTATTATCAGCTTACCGGAATTACACCAGCGAGGCCCTCCGCGGCCGTGAGGATGCCCTCA
    .....
AS:i:220    nn:i:0 tp:A:P cm:i:8 si:i:164    s2:i:0 def:f:0.0533   r1:i:0
NC_039553_1_4_430_13:0:0 11:0:0 _243d    163    NC_039553.1   4       60     150M =      281    427
    GGTTAGCTTAAACAAAGCATAACACTGCAAGATGATCCGCTGGCCGTGATAAGCCGAGCAGGAGTGAAGTTTGCTGAGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCCGCGGCCGTGAGGATGCCCTCA
    .....
AS:i:170    nn:i:0 tp:A:P cm:i:9 si:i:135    s2:i:0 def:f:0.0867   r1:i:0
NC_039553_1_5_495_6:0:0 11:0:0 _1d49    163    NC_039553.1   5       60     150M =      346    491
    GTGTAGCTTACACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTGATCAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGTTTCCCCAATTACACATGCGAGCCTCCGCGGCCGTGAGGATGCCCTCA
    .....
AS:i:240    nn:i:0 tp:A:P cm:i:12 si:i:148    s2:i:0 def:f:0.04     r1:i:0
NC_039553_1_6_523_6:0:0 9:0:0 82c      163    NC_039553.1   6       60     150M =      374    518
    TGTAGCTTAAACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTGATAAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTATTACAGCTTACCCAAATTACACATGCGAGCCTCCGCGGCCGTGAGGATGCCCTCA
    .....
AS:i:240    nn:i:0 tp:A:P cm:i:10 si:i:157    s2:i:0 def:f:0.04     r1:i:0
NC_039553_1_9_574_7:0:0 7:0:0 _181b    163    NC_039553.1   9       60     150M =      425    566
    AGCTTAAACAAAGCATAACACTGCAAGATGTTAAGCTGGCCGTGATAAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCCGCGGCCGTGAGGATGCCCTCC
    .....
AS:i:230    nn:i:0 tp:A:P cm:i:12 si:i:176    s2:i:0 def:f:0.0467   r1:i:0
NC_039553_1_10_391_9:0:0 8:0:0 256b    163    NC_039553.1   10      60     150M =      242    382
    GCTTAACAAAGCCTAACACTGAAAGATGTTAAGATGGCCGTGATAAGCCCCAACAGCACAGAAAGTTAGGTCTGGCTTATTACAGCTTACCCAAATTAGACATGCGAGCCTCCGCGGCCGTGATGCTGCCCTAGCCTCC
    .....
AS:i:210    nn:i:0 tp:A:P cm:i:5 si:i:156    s2:i:0 def:f:0.06     r1:i:0
NC_039553_1_11_509_6:0:0 11:0:0 _1i19    99     NC_039553.1   11      60     150M =      360    499
    CCTCACAAAGCATAACACTGAAAGATGTTAAGATGGCCGTATAAGCCCACAGCACTGAAAGTTAGGTCTGGCTTATTATGAGCTTACCCAAATTACACATGCGAGCCTCCGCGGCCGTGAGGATGCCCTCAGCCTCC
    .....
AS:i:242    nn:i:0 tp:A:P cm:i:10 si:i:150    s2:i:0 def:f:0.04     r1:i:0
NC_039553_1_12_427_9:0:0 9:0:0 _157    163    NC_039553.1   12      60     150M =      278    416
    TAAACAAAGCATAACACTGAAAGATTCAAGATGGCCGTAAAGCCCCAACAGCACTGAAAGTTAGGTCTGGCTTATTATCAGCTTACCCAAATTACACATGCGAGCCTCGGGGGCCCGTGAAGGATGCCCTCCGCTCCGT
    .....
AS:i:210    nn:i:0 tp:A:P cm:i:8 si:i:158    s2:i:0 def:f:0.06     r1:i:0
NM:i:8 ms:i:220
NM:i:13 ms:i:170
NM:i:6 ms:i:240
NM:i:6 ms:i:240
NM:i:7 ms:i:230
NM:i:9 ms:i:210
NM:i:6 ms:i:242
NM:i:9 ms:i:210
```

Figure: SAM File of Clean Reads

Data Preprocessing

Figure: SAM File of Chimeric Reads

Feature Extraction Pipeline

- BAM files were processed with a Python script to build a TSV feature matrix.
- Used Pysam for parsing alignments and NumPy for computation.

Feature Extraction Pipeline

- Focused on three features linked to PCR-induced chimeras:
 - ① **Supplementary Alignment (SA)**: Detects split alignments; counts and metrics extracted from SA tags
 - ② **K-mer Composition Difference**: Breakpoints inferred; left/right segments compared using cosine and JS metrics.
 - ③ **Microhomology**: Overlap at junction quantified (length + GC content) within a defined window.
- Pipeline design and outputs to be validated by experts.

Feature Extraction Pipeline

read_id	label	read_length	mean_basenr	ref_start	3'strand	mappq	cigar	has_sa	sa_count	num_segs_in_sa	diff	co_sa_min	dk_sa_max	d_sa_mean	sa_mean	sa_same	sa_opp	st_sa_max	nr_sa_mean	sa_min	nr_sa_mean	softclip_le	softclip_le	n_total	clippbreakpoint	kmer	cosi	kmer_js_d	microhom	microhom
NC_09950	0	150	13 NC_09950	3	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.9720	0.97143	1	0	
NC_09950	0	150	13 NC_09950	4	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98591	0.98571	1	0	
NC_09950	0	150	13 NC_09950	5	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.95887	0.95714	0	0	
NC_09950	0	150	13 NC_09950	6	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97183	0.97143	1	1	
NC_09950	0	150	13 NC_09950	7	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98644	0.98571	1	0	
NC_09950	0	150	13 NC_09950	10	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97296	0.97143	0	0	
NC_09950	0	150	13 NC_09950	11	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	1	1	
NC_09950	0	150	13 NC_09950	12	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98648	0.98571	1	1	
NC_09950	0	150	13 NC_09950	12	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.95680	0.95714	1	1	
NC_09950	0	150	13 NC_09950	14	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	0	0	
NC_09950	0	150	13 NC_09950	15	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98649	0.98571	1	0	
NC_09950	0	150	13 NC_09950	17	0	60	146M4S	0	0	1	0	0	0	0	0	0	0	0	0	0	4	4	146	0	0.5	0	0			
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98649	0.98571	3	0	
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	3	0	
NC_09950	0	150	13 NC_09950	18	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98629	0.98571	3	0	
NC_09950	0	150	13 NC_09950	19	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	3	0	
NC_09950	0	150	13 NC_09950	20	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97221	0.97143	0	0	
NC_09950	0	150	13 NC_09950	21	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0	
NC_09950	0	150	13 NC_09950	23	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98667	0.98571	0	0	
NC_09950	0	150	13 NC_09950	25	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98629	0.98571	0	0	
NC_09950	0	150	13 NC_09950	26	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.9863	0.98571	1	0	
NC_09950	0	150	13 NC_09950	32	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97238	0.97143	2	1	
NC_09950	0	150	13 NC_09950	34	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	1	1	0	0	
NC_09950	0	150	13 NC_09950	34	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	0	0	
NC_09950	0	150	13 NC_09950	35	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0	
NC_09950	0	150	13 NC_09950	36	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98648	0.98571	0	0	
NC_09950	0	150	13 NC_09950	38	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	1	0	
NC_09950	0	150	13 NC_09950	39	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98684	0.98571	0	0	
NC_09950	0	150	13 NC_09950	41	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.97296	0.97143	2	0.5	
NC_09950	0	150	13 NC_09950	43	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	75	0.98611	0.98571	0	0	

Figure: TSV Dataset showing Clean Reads

Feature Extraction Pipeline

1	read_id	label	read_id	mean	ref_nuc	ref_stal	strand	mpap	cigar	has_sai	sai_csq	num_sai	sai_dif	sai_min	sai_max	sai_maj	sai_pct	sai_maj	sai_min	sai_max	sai_pct	sai_pct	total_c	breakup	kmer_c	kmer_l	microf	microf	biology_gc			
19985	chimera_3	1	150	40	NC_03956	40	1	60	150M	0	0	0	1	0	0	0	0	0	0	0	0	0	0	75	0.98848	0.98571	0	0	0			
19986	chimera_3	1	150	40	NC_03956	53	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98848	0.98571	0	0	0			
19987	chimera_3	1	150	40	NC_03956	65	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98774	0.98571	0	0	0			
19988	chimera_3	1	150	40	NC_03956	66	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98774	0.98571	0	0	0			
19989	chimera_3	1	150	40	NC_03956	67	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98774	0.98571	0	0	0			
19990	chimera_3	1	150	40	NC_03956	67	1	60	110M32S	1	1	2	0	4246	4246	4246	0	1	10	10	0	0	0	32	32	32	118	1	1	0	0	0
19991	chimera_3	1	150	40	NC_03956	69	1	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98774	0.98571	0	0	0			
19992	chimera_3	1	150	40	NC_03956	76	0	60	100M41S	1	1	2	0	4237	4237	4237	0	1	16	16	0	0	0	41	41	41	109	1	1	0	0	0
19993	chimera_3	1	150	40	NC_03956	77	1	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.94396	0.94286	0	0	0			
19994	chimera_3	1	150	40	NC_03956	79	0	60	100M44S	1	1	2	0	4234	4234	4234	0	1	17	17	0	0	0	44	44	44	106	1	1	0	0	0
19995	chimera_3	1	150	40	NC_03956	84	0	60	120M38S	1	1	2	0	5197	5197	5197	0	1	10	10	0	0	0	38	38	38	101	0.98377	0.98443	0	0	0
19996	chimera_3	1	150	40	NC_03956	85	0	60	110M39S	1	1	2	0	5180	5180	5180	0	1	20	20	0	0	0	59	59	59	111	0.98394	0.98447	0	0	0
19997	chimera_3	1	150	40	NC_03956	88	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	75	0.98656	0.98554	1	1	1				
19998	chimera_3	1	150	40	NC_03956	89	0	60	150M39M	0	0	1	0	0	0	0	0	0	0	0	0	0	15	0	15	15	1	1	0			
19999	chimera_3	1	150	40	NC_03956	89	0	60	305120M	1	1	2	0	1973	1973	1973	0	1	1	1	0	0	0	30	0	30	30	0	0	0		
20000	chimera_3	1	150	40	NC_03956	89	0	60	415109M	1	1	2	0	1962	1962	1962	0	1	15	15	0	0	0	41	0	41	41	0.98411	0.98402	0	0	0
20001	chimera_3	1	150	40	NC_03956	89	1	60	969M54S	1	1	2	0	4224	4224	4224	0	1	48	48	0	0	0	54	54	96	1	1	0	0	0	
20002	chimera_3	1	150	40	NC_03956	89	1	60	335117M	1	1	2	0	1970	1970	1970	0	1	1	1	0	0	0	33	0	33	33	0.98275	0.98389	0	0	0
20003	chimera_3	1	150	40	NC_03956	90	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.95832	0.95714	1	1	1			
20004	chimera_3	1	150	40	NC_03956	90	1	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.95832	0.95714	1	1	1			
20005	chimera_3	1	150	40	NC_03956	90	1	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.95832	0.95714	1	1	1			
20006	chimera_3	1	150	40	NC_03956	91	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98774	0.98571	0	0	0			
20007	chimera_3	1	150	40	NC_03956	91	0	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.98332	0.98571	0	0	0			
20008	chimera_3	1	150	40	NC_03956	91	1	60	150M	0	0	1	0	0	0	0	0	0	0	0	0	0	0	75	0.95832	0.95714	0	0	0			
20009	chimera_3	1	150	40	NC_03956	91	1	60	949M65S	1	1	2	0	4222	4222	4222	0	1	52	52	0	0	0	56	56	94	1	1	0	0	0	
20010	chimera_3	1	150	40	NC_03956	92	0	60	295121M	0	0	1	0	0	0	0	0	0	0	0	0	0	29	29	0	29	0.98167	0.98338	1	0	0	
20011	chimera_3	1	150	40	NC_03956	92	0	60	275278M	1	1	2	0	3064	3064	3064	0	1	60	60	0	0	0	72	0	72	72	0.98608	0.98571	1	0	0
20012	chimera_3	1	150	40	NC_03956	92	0	60	66584M	1	1	2	0	3070	3070	3070	0	1	59	59	0	0	0	66	66	66	66	0.98611	0.98571	1	0	0
20013	chimera_3	1	150	40	NC_03956	92	0	60	100M41H	0	0	1	0	0	0	0	0	0	0	0	0	0	11	0	11	11	0	0	0			
20014	chimera_3	1	150	40	NC_03956	92	0	60	165124M	0	0	1	0	0	0	0	0	0	0	0	0	0	104	104	104	104	0	0	0			
20015	chimera_3	1	150	40	NC_03956	92	0	60	35147M	0	0	1	0	0	0	0	0	0	0	0	0	0	3	0	3	3	0	0.95	0.98041	1	0	0
20016	chimera_3	1	150	40	NC_03956	92	0	60	54996M	1	1	2	0	3082	3082	3082	0	1	30	30	0	0	0	54	0	54	54	0.9855	0.98534	1	0	0

Figure: TSV Dataset showing Chimeric Reads

Stratified Train–Test Split

- First step: **create a held-out test set** for final evaluation.
- Use `build_datasets.py`:
 - ① Combine clean and chimeric feature tables.
 - ② Attach labels (0 = clean, 1 = chimeric) if missing.
 - ③ Shuffle and perform **stratified** split:

Train : Test = 80% : 20%

with the same class proportions in each split.

- Output:
 - `train.tsv` (used for model selection and cross-validation).
 - `test.tsv` (kept untouched until the very end).

5-Fold Stratified Cross-Validation

- On the **training set only**, we perform:

5-fold stratified cross-validation

- Procedure:

- Split training data into 5 folds with balanced 0/1 labels.
- For each fold:
 - Train the model on 4 folds.
 - Evaluate on the remaining fold.
- Average metrics across the 5 folds:

mean F1 \pm std, mean accuracy \pm std

- This tells us:

- Typical performance** on unseen data.
- Stability** of each model (via standard deviation).
- Helps guide which algorithms are promising before going to the test set.

Model Zoo: Algorithms Compared

- We implemented a panel of 13 classifiers using scikit-learn and gradient boosting libraries:
 - **Baseline:** Dummy (always predicts most frequent class).
 - **Linear models:** Logistic regression (logreg_12), linear SVM with calibration.
 - **Tree ensembles:**
 - Random Forest, Extra Trees.
 - Gradient Boosting (sklearn).
 - XGBoost, LightGBM, CatBoost.
 - Bagging with decision trees.
 - **Others:** k-NN, Gaussian Naive Bayes, shallow MLP.
- All models use the same preprocessing pipeline:

Imputer (median) → StandardScaler → Classifier

Hyperparameter Tuning for Top Models

- For the 10 strongest families, we perform **RandomizedSearchCV** with 5-fold CV:
 - Logistic regression, linear SVM (calibrated).
 - Random Forest, Extra Trees, Gradient Boosting.
 - XGBoost, LightGBM, CatBoost.
 - Bagging (trees), MLP.
- Each search explores combinations of:
 - Tree depth, number of estimators, learning rate, subsample ratios, etc.
 - For MLP: hidden layer sizes, regularization (α), learning rate.
- Selection criterion:
 - Choose the hyperparameters with the best **cross-validated F1-score**.
 - Re-fit the best model on the **full training set**, then evaluate on the held-out test set.

Classification Metrics (Per-Read)

- For each model, on the test set we compute:

- Accuracy:**

$$\frac{\# \text{ correct predictions}}{\# \text{ all predictions}}$$

- Precision** (for chimeras):

$$\frac{TP}{TP + FP}$$

Of the reads we call “chimeric”, how many are truly chimeric?

- Recall** (for chimeras):

$$\frac{TP}{TP + FN}$$

Of all true chimeric reads, how many did we detect?

- F1-score** (for chimeras):

$$F1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Harmonic mean: high only if both precision and recall are high.

Threshold-Free Metrics: ROC–AUC and PR Curves

- Our models output a **score** per read (probability of being chimeric).
- By sweeping a threshold on this score, we can draw:
 - **ROC curve:**
 - x-axis: False Positive Rate (FPR).
 - y-axis: True Positive Rate (TPR = recall).
 - **ROC–AUC** = area under the curve.
 - **Precision–Recall (PR) curve:**
 - x-axis: Recall.
 - y-axis: Precision.
 - **Average Precision (AP)** = area under PR curve.
- Intuition for ROC–AUC:

$\text{AUC} \approx 0.84 \Rightarrow 84\% \text{ chance a random chimera is scored higher than a random non-chimera}$

Overall Performance Across Models (Test Set)

Model	CV Acc	CV F1	Test Acc	Test F1	ROC-AUC
Dummy baseline	0.50	0.67	0.50	0.67	0.50
Logistic regression	0.79	0.75	0.79	0.74	0.82
Linear SVM (cal.)	0.79	0.75	0.79	0.74	0.82
Random Forest	0.80	0.77	0.79	0.75	0.83
Extra Trees	0.80	0.77	0.79	0.75	0.82
Gradient Boosting	0.81	0.78	0.80	0.77	0.84
XGBoost	0.81	0.77	0.80	0.76	0.84
LightGBM	0.81	0.77	0.80	0.76	0.84
CatBoost	0.81	0.78	0.80	0.77	0.84
k-NN	0.78	0.75	0.78	0.75	0.81
Gaussian NB	0.75	0.66	0.74	0.65	0.82
Bagging (trees)	0.80	0.77	0.79	0.76	0.84
MLP	0.79	0.75	0.79	0.75	0.82

Table: Summary of cross-validation and test performance (chimeric class F1).

ROC and PR Curves (Placeholder)

ROC curves (CatBoost, GBM, RF,
logreg)

Precision–Recall curves

- ROC–AUC for top models: ≈ 0.84 .
- Curves pushed towards the top-left / top-right illustrate strong separation between clean and chimeric reads across thresholds.

Confusion Matrix and Class-Wise Behaviour (CatBoost)

Placeholder: confusion matrix for CatBoost on test set

CatBoost (test set, illustrative)

clean: precision ≈ 0.73 , recall ≈ 0.95

chimeric: precision ≈ 0.92 , recall ≈ 0.66

overall accuracy ≈ 0.80

- **Clean reads:**

- Very high recall: most true clean reads are correctly kept.

- **Chimeric reads:**

- High precision: when we call a read chimeric, it is usually correct.
- Moderate recall: we detect about two-thirds of all chimeras.

- Practical behaviour: **conservative chimera filter** that prioritizes not discarding clean reads.

Effect of Hyperparameter Tuning (F1 and ROC–AUC)

Model	F1 (base)	AUC (base)	F1 (tuned)	AUC (tuned)
CatBoost	0.767	0.839	0.769	0.844
Gradient Boosting	0.766	0.840	0.767	0.843
LightGBM	0.764	0.838	0.766	0.842
XGBoost	0.765	0.839	0.765	0.839
Random Forest	0.755	0.834	0.763	0.842
Bagging (trees)	0.760	0.837	0.763	0.842
Extra Trees	0.753	0.824	0.760	0.837
MLP	0.748	0.819	0.749	0.821
Logistic reg.	0.744	0.821	0.743	0.818
Linear SVM (cal.)	0.744	0.820	0.743	0.818

Table: Test F1 and ROC–AUC before vs after hyperparameter tuning.

- Tuning yields **modest but consistent gains** in F1 and ROC–AUC.
- Confirms that the initial defaults were already reasonable, but performance can be further refined.

Permutation Feature Importance (Placeholder)

- Top features across CatBoost, GBM, RF:

- total_clipped_bases
- kmer_js_divergence, kmer_cosine_diff
- softclip_left, softclip_right
- mapq

Placeholder: permutation importance for CatBoost

- Interpretation:

- Chimeras are characterized by **large clipped segments and abrupt k-mer composition shifts.**
- Aligners are already “seeing” the breakpoint signal; the ML model learns to combine these signals into a chimera score.



Summary of Findings

- We built a per-read feature table capturing:
 - Alignment and clipping patterns.
 - Supplementary alignments and breakpoint distances.
 - Sequence-level k-mer divergence and microhomology.
- A broad panel of ML models was evaluated:
 - Tree-based ensembles (CatBoost, Gradient Boosting, Random Forest, LightGBM, XGBoost) achieved the **best performance**.
 - Test F1 for chimeras $\approx 0.76\text{--}0.77$, ROC-AUC ≈ 0.84 .
- Model behaviour:
 - Conservative on clean reads (high recall).
 - High precision on chimeric reads, moderate recall.

Implications for Mitochondrial Assembly

- The ML classifier can be used as a **pre-filter** before assembling mitochondrial genomes:
 - Remove high-confidence chimeric reads to reduce false junctions.
 - Retain the majority of clean reads to preserve coverage.
- Especially useful for:
 - Small, circular, and repetitive organellar genomes where chimeras are particularly harmful.
 - Scenarios without high-quality reference genomes or abundance information.
- The feature importance analysis provides biological insight:
 - Confirms the role of soft-clipping, supplementary alignments, and k-mer jumps as core signals of chimeric structure.

Limitations and Future Work

- Current study uses **simulated** chimeras and a single species:
 - Need to validate on real experimental datasets.
 - Extend to other organellar genomes and library preparations.
- Classifier currently treats each read independently:
 - Future work: incorporate read-pair information, local read depth, or graph features.
- Integration into practical pipelines:
 - Wrap as a command-line tool interfacing with standard BAM/FASTQ workflows.
 - Benchmark impact on final assembly quality (contiguity, misassemblies).

Conclusion

- We developed a **machine learning pipeline** that:
 - Learns from alignment- and sequence-based features.
 - Achieves strong separation between clean and chimeric reads.
- Tree-based gradient boosting models (CatBoost, GBM, RF) provide:
 - High test F1 and ROC-AUC.
 - Interpretable feature importance aligned with known chimera mechanisms.
- This framework is a step towards **reference-free chimera detection** tailored for organellar genomes and low-resource settings.

Thank You

Questions?